Entropic Inference*

Ariel Caticha

Department of Physics, University at Albany-SUNY, Albany, NY 12222, USA.

Abstract

In this tutorial we review the essential arguments behing entropic inference. We focus on the epistemological notion of information and its relation to the Bayesian beliefs of rational agents. The problem of updating from a prior to a posterior probability distribution is tackled through an eliminative induction process that singles out the logarithmic relative entropy as the unique tool for inference. The resulting method of Maximum relative Entropy (ME), includes as special cases both MaxEnt and Bayes' rule, and therefore unifies the two themes of these workshops – the Maximum Entropy and the Bayesian methods – into a single general inference scheme.

1 Introduction

Our subject is inductive inference. Our goal in this tutorial paper is to review the problem of updating from a prior probability distribution to a posterior distribution when new information becomes available.

First we tackle the question of the nature of information itself: What is information? It is clear that data "contains" or "conveys" information, but what does this precisely mean? Is information physical? We discuss how in a properly Bayesian framework one can usefully adopt a concept of information that is more directly related to the epistemological concerns of rational agents.

Then we turn to the actual methods to process information. We argue for the uniqueness and universality of the Method of Maximum relative Entropy (ME) and then we discuss its relation to Bayesian methods. At first sight Bayesian and Maximum Entropy methods appear unrelated. Bayes' rule is the natural way to update probabilities when the new information is in the form of data. On the other hand, Jaynes' method of maximum entropy, MaxEnt, is designed to handle information in the form of constraints [1]. An important question is whether they are compatible with each other. We show that the ME method includes both MaxEnt and Bayesian methods as special cases and

^{*}Presented at MaxEnt 2010, the 30th International Workshop on Bayesian Inference and Maximum Entropy Methods in Science and Engineering (July 4-9, 2010, Chamonix, France).

allows us to extend them to situations that lie beyond the reach of either of them individually.

Finally we explore an important extension of the ME method. The distribution of maximum entropy has the highest probability of being the correct choice of posterior, but how justified are we in ruling out those distributions that do not maximize the entropy? The extended ME assigns a probability to those other distributions and this has a wide variety of applications: it provides a connection to the theory of large deviations, to fluctuation theory, to entropic priors, and most recently to quantum mechanics. The possibilities are endless.

We make no attempt to provide a review of the literature on entropic inference. The following list, which reflects only some contributions that are directly related to the particular approach described in this tutorial, is incomplete but might nevertheless be useful: Jaynes [1], Shore and Johnson [2], Williams [3], Skilling [4], Rodríguez [5][6], Giffin and Caticha [7]-[12].

2 What is information?

The expression that systems "carry" or "contain" information can perhaps be traced to Shannon's theory of communication: a system is analogous to a message. The system "carries" information about its own state and, in this sense, one can say that *information is physical*. Such physical information is directly associated to the system. Our interest here is in an altogether different notion of information which we might call epistemological and which is directly associated to the beliefs of rational agents. Indeed, any fully Bayesian theory of information requires an explicit account of how such epistemological information is related to rational beliefs.

The need to update from one state of belief to another is driven by the conviction that not all probability assignments are equally good; some beliefs are preferable to others in the very pragmatic sense that they enhance our chances to successfully navigate this world. The idea is that, to the extent that we wish to be called rational, we will improve our beliefs by revising them when new information becomes available: *Information is what forces a change of rational beliefs*. Or, to put it more explicitly: *Information is a constraint on rational beliefs*.

This definition – information is a constraint – is sufficient for our present purposes but would benefit from further elaboration. The definition captures a notion of information that is directly related to changing our minds: information is the driving force behind the process of learning. It incorporates an important feature of rationality: being rational means accepting that our beliefs must be constrained in very specific ways – not everything goes. But surely this is not enough: the indiscriminate acceptance of any arbitrary constraint does not qualify as rational behavior. To be rational an agent must exercise some judgement before accepting a particular piece of information as a reliable basis for the revision of its beliefs and this raises questions about what judgements might be considered sound. Indeed, there is no implication that the information

must be true; only that we accept it as true. False information is information too, at least as long as we are prepared to accept it and allow it to affect our beliefs.

The paramount virtue of our definition is that it is useful. It allows precise quantitative calculations even though the notion of an amount of information, whether measured in bits or otherwise, is not introduced. By information in its most general form, we just mean the set of constraints on the family of acceptable posterior distributions and this is precisely the kind of information the method of maximum entropy is designed to handle.

3 Updating probabilities: the ME method

The uncertainty about a variable $x \in \mathcal{X}$ (whether discrete or continuous, in one or several dimensions) is described by a probability distribution q(x). Our goal is to design a method to update from a prior distribution q(x) to a posterior distribution P(x) when new information in the form of constraints becomes available. (The constraints can be given in terms of expected values but this is not necessary. Other types of constraints are allowed too; an example is appears in section 5.)

The problem is to select a distribution from among all those that satisfy the constraints. The procedure is to rank the candidate distributions in order of increasing preference [4]. It is clear that to accomplish our goal the ranking must be transitive: if distribution p_1 is preferred over p_2 , and p_2 is preferred over p_3 , then p_1 is preferred over p_3 . Such transitive rankings are implemented by assigning to each p(x) a real number S[p] in such a way that if p_1 is preferred over p_2 , then $S[p_1] > S[p_2]$. The selected distribution P (one or possibly many, for on the basis of the available information we might have several equally preferred distributions) will be that which maximizes the quantity S[p], which we will henceforth call entropy. We are thus led to a method of Maximum Entropy (ME) that involves entropies that are real numbers and that are meant to be maximized. These features are imposed by design; they are dictated by the function that the ME method is being designed to perform and not by any objective properties of the external world.

Next we must make a definite choice for the functional S[p]. Since the purpose of the method is to update from priors to posteriors the ranking scheme must depend on the particular prior q and therefore the entropy S must be a functional of both p and q. Thus the entropy S[p,q] produces a ranking of the distributions p relative to the given prior q: S[p,q] is the entropy of p relative to q. Accordingly S[p,q] is commonly called relative entropy, but since all entropies are relative, even when relative to a uniform distribution, the modifier 'relative' is redundant and can be dropped.

The functional S[p,q] is selected by a process of *eliminative induction*. The idea is simple: we start with a *sufficiently broad* family of candidates and identify a number of special cases for which we know what the preferred distribution ought to be. Then we just eliminate all those candidates that fail to provide the

right update. As we shall see the selection criteria adopted below are sufficiently constraining that there is a single entropy functional S[p,q] that survives the process of elimination.

This approach has a number of virtues. First, to the extent that the selection criteria are universally desirable, then the single surviving entropy functional will be of universal applicability too. Second, the reason why any entropy candidate is eliminated is quite explicit – at least one of the selection criteria is violated. Thus, the justification behind the single surviving entropy is not that it leads to demonstrably correct inferences, but rather, that other entropies are demonstrably wrong.

The selection criteria are chosen to reflect the conviction that information collected in the past and codified into the prior distribution is valuable and should not be ignored. This attitude is very conservative: the only aspects of one's beliefs that should be updated are those for which new evidence has been supplied. Moreover, as we shall see below, the selection criteria merely tell us what not to update, which has the virtue of maximizing objectivity – there are many ways to change something but only one way to keep it the same. These ideas are summarized in the following

Principle of Minimal Updating (PMU): Beliefs must be revised only to the extent required by the new information.

Three selection criteria, a brief motivation for them, and their consequences for the functional form of the entropy are listed below (proofs and more details are given in [12]). The reason these criteria are so constraining is that they refer to three infinitely large classes of special cases where the desired update is known.

Criterion 1: Locality. Local information has local effects.

If the information to be processed does not refer to an x in a particular subdomain $\mathcal{D} \subset \mathcal{X}$ then the PMU requires that we do not change our minds about $x \in \mathcal{D}$. More precisely, we require that the prior conditioned on \mathcal{D} is not updated. The selected posterior is such that $P(x|\mathcal{D}) = q(x|\mathcal{D})$. Dropping additive terms and multiplicative factors that do not affect the overall ranking, the surviving entropy functionals are of the form

$$S[p,q] = \int dx F(p(x), q(x), x) , \qquad (1)$$

where F is some unknown function and by $\int dx$ we mean a discrete sum or continuous integral (possibly over several dimensions) as the case might require.

Criterion 2: Coordinate invariance. The system of coordinates carries no information.

The points x can be labeled in different ways using different coordinate systems but this should not affect the ranking of the distributions. The consequence of criterion 2 is that the surviving entropies can be written as

$$S[p,q] = \int dx \, m(x) \Phi\left(\frac{p(x)}{m(x)}, \frac{q(x)}{m(x)}\right) , \qquad (2)$$

where m(x) is a probability density, which implies that dx m(x), p(x)/m(x), and q(x)/m(x) are coordinate invariants. (Again, additive terms and multiplicative factors that do not affect the overall ranking have been dropped.) We see that the single unknown function F in (1) with three arguments has been replaced by two unknown functions. One is the density m(x), and the other is a function Φ with two arguments. The density m(x) is determined by invoking the locality criterion once again.

Criterion 1 (a special case): When there is no new information there is no reason to change one's mind.

When no new information is available the domain \mathcal{D} in criterion 1 coincides with the whole space \mathcal{X} . The conditional probabilities $q(x|\mathcal{D}) = q(x|\mathcal{X}) = q(x)$ should not be updated and the selected posterior coincides with the prior, P(x) = q(x). The consequence is that up to normalization the unknown m(x) must be the prior distribution q(x). The entropy is now restricted to functionals of the form

$$S[p,q] = \int dx \, q(x) \Phi\left(\frac{p(x)}{q(x)}\right) . \tag{3}$$

Criterion 3: Independence. When systems are known to be independent it should not matter whether they are treated separately or jointly.

The preservation of independence is a particularly important concern for science because without it science is not possible. The reason is that in any inference it is assumed that the universe is partitioned into the system of interest and other systems that constitute the rest of the universe. What is important about those other systems is precisely that they can be ignored – whether they are included in the analysis or not should make no difference. If they did matter they should have been incorporated as part of the system of interest in the first place.

It is crucial that Criterion 3 be applied to all independent systems whether they are identical or not, whether just two or many, or even infinitely many. This criterion is sufficiently constraining that (up to additive terms and multiplicative factors that do not affect the overall ranking scheme) there is a single surviving entropy functional given by the usual logarithmic relative entropy [12],

$$S[p,q] = -\int dx \, p(x) \log \frac{p(x)}{q(x)} \tag{4}$$

These results are summarized as follows:

The ME method: The objective is to update from a prior distribution q to a posterior distribution P given the information that the posterior lies within a certain family of distributions p. The selected posterior P is that which maximizes the entropy S[p,q]. Since prior information is valuable the functional S[p,q] is chosen so that beliefs are updated only to the minimal extent required by the new information. No interpretation for S[p,q] is given and none is needed.

4 Bayes' rule and its generalizations

Bayes' rule is used to make inferences about one or several quantities $\theta \in \Theta$ on the basis of information in the form of data $x \in \mathcal{X}$. More specifically, the problem is to update our beliefs about θ on the basis of three pieces of information: (1) the prior information codified into a prior distribution $q(\theta)$; (2) the data $x \in \mathcal{X}$ (obtained in one or many experiments); and (3) the known relation between θ and x given by the model as defined by the sampling distribution or likelihood function, $q(x|\theta)$. The updating consists of replacing the *prior* probability distribution $q(\theta)$ by a *posterior* distribution $P(\theta)$ that applies after the data has been processed.

Remark: We emphasize that the information about how x is related to θ is contained in the *functional form* of the distribution $q(\cdot|\theta)$ which is completely unrelated to the actual values of the observed data.

The insight that will allow Bayes' rule to be smoothly incorporated into the entropic inference framework [3][9] is that the relevant universe of discourse is not Θ but the product space $\Theta \times \mathcal{X}$ [5][6]. We deal with joint distributions and the relevant *joint prior* is $q(x,\theta) = q(\theta)q(x|\theta)$.

Remark: Bayes' rule is usually written in the form

$$q(\theta|x) = q(\theta) \frac{q(x|\theta)}{q(x)} , \qquad (5)$$

and called Bayes' theorem. This formula is a restatement of the product rule. It is valid for any value of x whether it coincides with the observed data or not and therefore it is a simple consequence of the internal consistency of the prior beliefs. Within the framework of entropic inference the left hand side is not a posterior but rather a prior probability – it is the prior probability of θ conditional on x.

Next we collect data and the observed values turn out to be x'. This constrains the posterior to the family of distributions $p(x, \theta)$ defined by

$$p(x) = \int d\theta \, p(\theta, x) = \delta(x - x') \ . \tag{6}$$

This data information is not, however, sufficient to determine the joint distribution

$$p(x,\theta) = p(x)p(\theta|x) = \delta(x - x')p(\theta|x'). \tag{7}$$

Any choice of $p(\theta|x')$ is in principle possible. Within the framework of entropic inference (see [9]) the joint posterior $P(x,\theta)$ is the minimal update from the prior $q(x,\theta)$ that agrees with the data constraint. To find it maximize the entropy,

$$S[p,q] = -\int dx d\theta \ p(x,\theta) \log \frac{p(x,\theta)}{q(x,\theta)} , \qquad (8)$$

subject to the *infinite* number of constraints given by eq. (6). Note that there is one constraint for each value of x. The corresponding Lagrange multipliers are denoted $\lambda(x)$. Maximizing (8) subject to (6) and normalization,

$$\delta \left\{ S + \alpha \left[\int dx d\theta \ p(x,\theta) - 1 \right] + \int dx \ \lambda(x) \left[\int d\theta \ p(x,\theta) - \delta(x - x') \right] \right\} = 0 \ , \ (9)$$

yields

$$P(x,\theta) = q(x,\theta) \frac{e^{\lambda(x)}}{Z} , \qquad (10)$$

where Z is a normalization constant, and $\lambda(x)$ is determined from (6),

$$\int d\theta \ q(x,\theta) \frac{e^{\lambda(x)}}{Z} = q(x) \frac{e^{\lambda(x)}}{Z} = \delta(x - x') \ , \tag{11}$$

so that the joint posterior is

$$P(x,\theta) = q(x,\theta) \frac{\delta(x-x')}{q(x)} = \delta(x-x')q(\theta|x) . \tag{12}$$

The corresponding marginal posterior probability $P(\theta)$ is

$$P(\theta) = \int dx \, P(\theta, x) = q(\theta | x') = q(\theta) \frac{q(x'|\theta)}{q(x')} , \qquad (13)$$

which coincides with Bayes' rule. This is intuitively reasonable: we maintain those beliefs about θ that are consistent with the data values x' that turned out to be true. Data values that were not observed are discarded because they are now known to be false. The extension to repeatable independent experiments is straightforward [12].

Next I give a couple of very simple examples that show how entropic methods allow generalizations of Bayes' rule.

Example 1.– Jeffrey's rule. As before, the prior information consists of our prior knowledge about θ given by the distribution $q(\theta)$ and the relation between x and θ is given by the likelihood $q(x|\theta)$. But now the information about x is limited because the data is uncertain. The marginal posterior p(x) is no longer a sharp delta function but some other known distribution, $p(x) = P_D(x)$. This is still an infinite number of constraints

$$p(x) = \int d\theta \, p(\theta, x) = P_D(x) . \tag{14}$$

Maximizing (8) subject to (14) and normalization, leads to

$$P(x,\theta) = P_D(x)q(\theta|x) . (15)$$

The corresponding marginal posterior,

$$P(\theta) = \int dx \, P_D(x) q(\theta|x) = q(\theta) \int dx \, P_D(x) \frac{q(x|\theta)}{q(x)} , \qquad (16)$$

is known as Jeffrey's rule. In the limit when the data are sharply determined $P_D(x) = \delta(x - x')$ the posterior reproduces Bayes' rule (13).

Example 2.– Unknown likelihood. The following example derives and generalizes Zellner's Bayesian Method of Moments [13]. Usually the relation between x and θ is given by a known likelihood function $q(x|\theta)$ but suppose this relation is not known. This is the case when the joint prior is so ignorant that

information about x tells us nothing about θ and vise versa; such a prior treats x and θ as statistically independent, $q(x,\theta) = q(x)q(\theta)$. Since we have no likelihood function the information about the relation between θ and the data x must be supplied elsewhere. One possibility is through a constraint. Suppose that in addition to normalization and the uncertain data constraint, eq.(14), we also know that the expected value over θ of a function $f(x,\theta)$ is

$$\langle f \rangle_x = \int d\theta \, p(\theta|x) f(x,\theta) = F(x) \ .$$
 (17)

We seek a posterior $P(x, \theta)$ that maximizes (8). Introducing Lagrange multipliers α , $\lambda(x)$, and $\gamma(x)$,

$$0 = \delta \left\{ S + \alpha \left[\int dx d\theta \ p(x,\theta) - 1 \right] + \int dx \ \lambda(x) \left[\int d\theta \ p(x,\theta) - P_D(x) \right] (18) \right. \\ \left. + \int dx \ \gamma(x) \left[\int d\theta \ p(x,\theta) f(x,\theta) - P_D(x) F(x) \right] \right\} , \tag{19}$$

the variation over $p(x, \theta)$ yields

$$P(x,\theta) = \frac{1}{\zeta} q(x) q(\theta) e^{\lambda(x) + \gamma(x) f(x,\theta)} , \qquad (20)$$

where ζ is a normalization constant. The multiplier $\lambda(x)$ is determined from (6),

$$P(x) = \int d\theta \, P(\theta, x) = \frac{1}{\zeta} q(x) e^{\lambda(x)} \int d\theta \, q(\theta) \, e^{\gamma(x)f(x,\theta)} = P_D(x)$$
 (21)

then,

$$P(x,\theta) = P_D(x) \frac{q(\theta) e^{\gamma(x)f(x,\theta)}}{\int d\theta' q(\theta') e^{\gamma(x)f(x,\theta')}}$$
(22)

so that

$$P(\theta|x) = \frac{P(x,\theta)}{P(x)} = \frac{q(\theta) e^{\gamma(x)f(x,\theta)}}{Z(x)} \quad \text{with} \quad Z(x) = \int d\theta' \, q(\theta') \, e^{\gamma(x)f(x,\theta')} \tag{23}$$

The multiplier $\gamma(x)$ is determined from (17)

$$\frac{1}{Z(x)} \frac{\partial Z(x)}{\partial \gamma(x)} = F(x) . {24}$$

The corresponding marginal posterior is

$$P(\theta) = \int dx \, P_D(x) P(\theta|x) = q(\theta) \int dx \, P_D(x) \frac{e^{\gamma(x)f(x,\theta)}}{Z(x)} . \tag{25}$$

In the limit when the data are sharply determined $P_D(x) = \delta(x - x')$ the posterior takes the form of Bayes theorem,

$$P(\theta) = q(\theta) \frac{e^{\gamma(x')f(x',\theta)}}{Z(x')} , \qquad (26)$$

where up to a normalization factor $e^{\gamma(x')f(x',\theta)}$ plays the role of the likelihood and the normalization constant Z plays the role of the evidence.

In conclusion, these examples demonstrate that the method of maximum entropy can fully reproduce the results obtained by the standard Bayesian methods and allows us to extend them to situations that lie beyond their reach such as when the likelihood function is not known. Other such examples are given in [11] and [12].

5 Deviations from maximum entropy

To complete the design of the ME method we must address one last issue. Once we have decided that the distribution that maximizes entropy is to be preferred over all others we ask: to what extent are the other distributions ruled out? The discussion below follows [8][12].

The original problem was to update from a prior q(x) given constraints that define the space Θ of acceptable distributions. We assume that these distributions, that is, the "points" in the space Θ , can be labelled by coordinates θ . Thus, Θ is a statistical manifold and its points can be written as $p(x|\theta)$. Maximizing S[p,q] over all the $p(x|\theta)$ in Θ leads to the preferred distribution, say $p(x|\theta_0)$.

The question about the extent that distributions with $\theta \neq \theta_0$ are ruled out is a question about the probability of various values of θ : to what extent do we believe that the selected value should lie within any particular range $d\theta$? Thus we are not just concerned with the probability of x, but with the joint distribution $p(x,\theta)$. To assign $p(x,\theta)$ we apply the same ME method but in the larger joint space: maximize the joint entropy

$$S[p,q] = -\int dx \, d\theta \, p(x,\theta) \log \frac{p(x,\theta)}{q(x,\theta)} \,, \tag{27}$$

for a suitable prior $q(x,\theta)$ and under the appropriate constraints.

Choosing a prior is always tricky because it represents what we knew before the relevant new information became available. We want to represent a state of extreme ignorance: the precise relation between θ s and xs is not (yet) known and therefore $q(x,\theta)$ is a product, $q(x,\theta) = q(x)q(\theta)$, so that knowing x tells us nothing about θ and vice versa. For q(x) we retain the prior used in the original problem where we updated from q(x) to $p(x|\theta_0)$.

For $q(\theta)$ we plead ignorance once again and choose a uniform distribution. This is somewhat trickier than may seem at first sight because *uniform* does not mean *constant*. The uniform distribution assigns equal probabilities to equal volumes in Θ and does not depend on the particular choice of coordinates. (A constant distribution, on the other hand, depends on the choice of coordinates: a distribution that is constant in one frame coordinate will not be constant in another.) This requires a well-defined notion of volume. Fortunately, the statistical manifold Θ is a metric space: there is a single unique geometry that properly takes into account the fact that the points in Θ are not structureless

points but are actual probability distributions. This is given by the Fisher-Rao information metric, $g_{ij}(\theta)$ [14][12]. The corresponding volume elements are given by $g^{1/2}(\theta)d^n\theta$, where $g = \det g_{ij}$. Therefore the uniform (unnormalized) prior is $q(\theta) = g^{1/2}(\theta)$ and the joint prior is $q(x,\theta) = g^{1/2}(\theta)q(x)$.

The crucial constraint on the joint distributions $p(x, \theta) = p(\theta)p(x|\theta)$ specifies the conditional distributions $p(x|\theta)$. This amounts to selecting the particular space Θ under consideration.

The preferred joint distribution $P(x, \theta)$ is that which maximizes the joint entropy S[p, q] over all normalized distributions of the form $p(x, \theta) = p(\theta)p(x|\theta)$ where we vary with respect to $p(\theta)$ and restrict to $p(x|\theta) \in \Theta$. It is convenient to rewrite (27) as

$$S[p,q] = -\int d\theta \, p(\theta) \log \frac{p(\theta)}{q^{1/2}(\theta)} + \int d\theta \, p(\theta) S(\theta), \tag{28}$$

where

$$S(\theta) = -\int dx \, p(x|\theta) \log \frac{p(x|\theta)}{q(x)}.$$
 (29)

The result is the probability that θ lies within a small volume $g^{1/2}(\theta)d^n\theta$,

$$P(\theta)d^n\theta = \frac{1}{\zeta} e^{S(\theta)} g^{1/2}(\theta) d^n\theta \quad \text{with} \quad \zeta = \int d^n\theta g^{1/2}(\theta) e^{S(\theta)}. \tag{30}$$

The preferred value of θ is that θ_0 which maximizes the entropy $S(\theta)$, eq.(29), because this maximizes the scalar probability density $\exp S(\theta)$. But it also tells us the degree to which values of θ away from the maximum are ruled out.

One of the limitations of the standard MaxEnt method is that it selects a single "posterior" $p(x|\theta_0)$ and all other distributions are strictly ruled out. The result (30) overcomes this limitation and finds many applications. For example, it extends the Einstein theory of thermodynamic fluctuations beyond the regime of small fluctuations; it provides a bridge to the theory of large deviations; and, suitably adapted for Bayesian data analysis, it leads to the notion of entropic priors.

6 Conclusions

Any Bayesian account of the notion of information cannot ignore the fact that Bayesians are concerned with the beliefs of rational agents. The relation between information and beliefs must be clearly spelled out. The definition we have proposed – that information is that which constrains rational beliefs and therefore forces the agent to change its mind – is convenient for two reasons. First, the information/belief relation is explicit, and second, the definition is ideally suited for quantitative manipulation using the ME method.

The main conclusion is that the logarithmic relative entropy is the only candidate for a general method for updating probabilities – the ME method – and this includes both MaxEnt and Bayes' rule as special cases; it unifies

them into a single theory of inductive inference and allows new applications. Indeed, much as the old MaxEnt method provided the foundation for statistical mechanics, recent work suggests that the extended ME method provides an entropic foundation for quantum mechanics.

Acknowledgements: I would like to acknowledge valuable discussions with C. Cafaro, N. Caticha, A. Giffin, K. Knuth, C. Rodríguez, J. Skilling, and C.-Y. Tseng.

References

- E. T. Jaynes, Phys. Rev. 106, 620 and 108, 171 (1957); R. D. Rosenkrantz (ed.), E. T. Jaynes: Papers on Probability, Statistics and Statistical Physics (Reidel, Dordrecht, 1983); E. T. Jaynes, Probability Theory: The Logic of Science (Cambridge University Press, Cambridge, 2003).
- [2] J. E. Shore and R. W. Johnson, IEEE Trans. Inf. Theory IT-26, 26 (1980);IEEE Trans. Inf. Theory IT-27, 26 (1981).
- [3] P. M. Williams, Brit. J. Phil. Sci. 31, 131 (1980).
- [4] J. Skilling, "The Axioms of Maximum Entropy" in Maximum-Entropy and Bayesian Methods in Science and Engineering, G. J. Erickson and C. R. Smith (eds.) (Kluwer, Dordrecht, 1988).
- [5] C. C. Rodríguez, "Entropic priors" in Maximum Entropy and Bayesian Methods, edited by W. T. Grandy Jr. and L. H. Schick (Kluwer, Dordrecht, 1991).
- [6] C. C. Rodríguez, "Are we cruising a hypothesis space?" in *Maximum Entropy and Bayesian Methods*, edited by W. von der Linden, V. Dose, R. Fischer and R. Preuss (Kluwer, Dordrecht, 1999).
- [7] A. Caticha, 'Maximum entropy, fluctuations and priors', in Bayesian Methods and Maximum Entropy in Science and Engineering, ed. by A. Mohammad-Djafari, AIP Conf. Proc. 568, 94 (2001) (arXiv.org/abs/math-ph/0008017).
- [8] A. Caticha, "Relative Entropy and Inductive Inference," in Bayesian Inference and Maximum Entropy Methods in Science and Engineering, ed. by G. Erickson and Y. Zhai, AIP Conf. Proc. 707, 75 (2004) (arXiv.org/abs/physics/0311093).
- [9] A. Caticha and A. Giffin, "Updating Probabilities", Bayesian Inference and Maximum Entropy Methods in Science and Engineering, Ali Mohammad-Djafari (ed.), AIP Conf. Proc. 872, 31 (2006) (arxiv.org/abs/physics/0608185).

- [10] A. Caticha, "Information and Entropy," in *Bayesian Inference and Maximum Entropy Methods in Science and Engineering*, ed. by K. Knuth *et al.*, AIP Conf. Proc. **954**, 11 (2007) (arXiv.org/abs/0710.1068).
- [11] A. Giffin and A. Caticha, "Updating Probabilities with Data and Moments," in *Bayesian Inference and Maximum Entropy Methods in Science and Engineering*, ed. by K. Knuth *et al.*, AIP Conf. Proc. **954**, 74 (2007) (arXiv.org/abs/0708.1593).
- [12] A. Caticha, Lectures on Probability, Entropy, and Statistical Physics (Max-Ent 2008, São Paulo, Brazil) (arXiv.org/0808.0012).
- [13] A. Zellner, "The Bayesian Method of Moments", Advances in Econometrics 12, 85 (1997).
- [14] S. Amari, Differential-Geometrical Methods in Statistics (Springer-Verlag, 1985).